

STAT 340: Lecture 1 - Introduction

Aliaksandr Hubin

February 10 2025

Applied Methods in Statistics

Lecturer: Ass. Prof. Aliaksandr Hubin, KBM

Assistant Teacher: Elias Bye Rossetnes

Date: 03.02.2025

Institution: NMBU - Norges miljø- og biovitenskapelige universitet

Course Teachers

Lecturer

Ass. Prof. Aliaksandr Hubin, KBM - Office: M252 - Old diary science building (Meieribygningen)

► Email: aliaksandr.hubin@nmbu.no

Assistant Teacher

Elias Bye Rossetnes

Class Schedule

Lectures

- ▶ **When:** Tuesdays 14:15–16:00
- ▶ **Where:** Room T330
- ▶ **First Lecture:** Feb 3

Exercises (individual work)

- ▶ **When:** Thursdays 14:15–16:00
- ▶ **Where:** Room BT3A10
- ▶ **First Session:** Feb 5

Exercises (solutions presentation)

- ▶ **When:** Mondays 10:15–12:00
- ▶ **Where:** Room T132
- ▶ **First Session:** Feb 9

Practical Information

Mandatory Activity

- ▶ **Four compulsory quiz assignments**
- ▶ You can be released from a quiz by answering 10 questions correctly during lectures (no penalty for wrong answers).
- ▶ This will work even after failing a quiz.

Exam

- ▶ **Format:** Multiple choice (100%)
- ▶ **Grading:** A–E or Not Passed
- ▶ **Allowed:** All calculators and specified **written** aids
- ▶ **Date & Time:** **May 18, 09:00–12:00**
- ▶ **Important:** No date changes possible

Learning Activities (Weekly)

- ▶ **Lectures:** 2 hours
- ▶ **Group Exercises:** 2 hours
- ▶ **Individual Work:** 2 hours

Total: 6 hours per week

Assumed Background: STAT 100

Students should be familiar with:

- ▶ Basic rules of probability (we shall repeat that today)
- ▶ Common probability distributions (we shall repeat that today):
 - ▶ Binomial
 - ▶ Gaussian (Normal)
 - ▶ Student's t
 - ▶ Fisher's F
- ▶ Hypothesis testing:
 - ▶ One-sample t -tests
 - ▶ Two-sample t -tests
 - ▶ Paired sample t -tests
- ▶ One-way Analysis of Variance (ANOVA)
- ▶ Simple linear regression (one predictor variable)
- ▶ Contingency tables

Software & Data

Software

- ▶ **R Studio** for R programming language
- ▶ Practical analysis using various statistical methods
- ▶ Theoretical and practical interpretation of results
- ▶ (Other software may be used if methods are available)

Data

- ▶ Mainly R library: **BIAS.data** (GitHub)
<https://github.com/thoree/BIAS.data>

STAT 340 Course Content

Main Topics

1. **Introduction to R and RStudio**

- ▶ Practical analysis using various statistical methods

2. **Multiple Linear Regression**

3. **Multivariate Analysis**

- ▶ Principal Component Analysis (PCA)

4. **Classification Methods**

5. **Clustering Methods**

6. **Generalized Linear Models (GLM)**

- ▶ Binomial/Bernoulli (logistic) regression
- ▶ Poisson regression
- ▶ Mixed models (GLMM): Random effects, repeated measurements

Required Learning Resources

Primary Materials

- ▶ **Lecture presentations** (provided)
- ▶ **All exercises** (provided)

Recommended Textbooks (available online)

Main Reference: - R Book: Michael J Crawley, Wiley (2nd ed.)

- Chapters: 1–9 (as needed), 10–13, 17, 19, 25 - Freely available

https:

[//www.cs.upc.edu/~robert/teaching/estadistica/TheRBook.pdf](https://www.cs.upc.edu/~robert/teaching/estadistica/TheRBook.pdf)

Supplementary (mathematical details): - ISL: James et al. “An Introduction to Statistical Learning” - Chapters: 1–4, 6.1, 7.1, 10

- Freely available <https://www.statlearning.com>

Typical Week Structure

ws	Monday 2.2	Tuesday 3.2	Wednesday 4.2	Thursday 5.2	Friday 6.2
8					
9					
10	11:00				
11	STAT340 T132 Hjelpørelser - Exercise				
12	13:00				
13					
14		15:00		14:00	
15		STAT340 T1330 Aleksandr Hubin Lecture		STAT340 B13010 Hjelpørelser - Exercise	
16		16:30		16:00	
17					

		Response				
		Continuous		Discrete		
		Normal	Non-normal	Counts	Bivariate	Multivariate
"Study of associations"		"Linear"	(Transform?)	Poisson	Logistic	Multinomial
Predictors	Fixed	Contin.	Linear Regression	GLM regression	GLM Regression	GLM regression or Discriminant analysis (LDA, QDA ++)
		Discrete	ANOVA	GLM ANOVA	GLM ANOVA	GLM ANOVA/Contingency tables or Discriminant analysis
		Both	ANCOVA Dummy regr	GLM ANCOVA	GLM ANCOVA	GLM ANCOVA or Discriminant analysis
	Random	Any kind	LMM	GLMM		
New sample		Any kind	Prediction			Classification

Reflection Activity

Turn to your neighbour(s) and discuss:

- ▶ Do you know what kind of **data** you will be dealing with in your work?
- ▶ Do you have any idea what **association studies** or **methods** are relevant for your research?
- ▶ Are they mentioned in the table in the previous page?

Now, Let's Begin!

Lecture 1, Part B: Review of Basic Statistics

We'll start with basics of **probability theory** - the mathematical foundation for **uncertainty** and **statistical inference**.

Part 0: Foundations of Probability

Setting Up a Random Experiment

Before studying random variables, we need to define randomness formally.

A probability setup has three components:

1. **Sample space** Ω : all possible outcomes
 - ▶ *Discrete*: Coin flip: $\Omega = \{H, T\}$
 - ▶ *Continuous*: Adult height: $\Omega = [120, 250]$ cm
2. **Events** $\mathcal{F}$: subsets of outcomes we observe
 - ▶ *Discrete*: $\{H\} = \text{"heads"}$
 - ▶ *Continuous*: $\{h : 175.3\}$
3. **Probabilities** P : assigns to events a number in $[0, 1]$
 - ▶ *Discrete*: $P(H) = 0.5$
 - ▶ *Continuous*: $P(160 \leq h \leq 180) = 0.68$

Together: $(\Omega, \mathcal{F}, P)$ is a **probability space**.

Andrey Kolmogorov (1903–1987)



Figure 1: Soviet mathematician who revolutionized probability by providing rigorous mathematical foundations

- ▶ Established the **foundations of probability** (1933)
- ▶ Formalized the concept of probability space $(\Omega, \mathcal{F}, P)$.

Kolmogorov's Axioms:


1. $P(\Omega) = 1$ (certainty has probability 1)
2. $P(A) \geq 0$ for all outcomes A (non-negativity)
3. Countable additivity: If $A_1, A_2, \dots$ are disjoint events, then
$$P(\cup_i A_i) = \sum_i P(A_i)$$

Part 1: Random Variables

What is a Random Variable?

A **random variable** is a function that assigns a numeric value to each outcome in the sample space. Maps from experiment sample space to a real space.

Why? Numbers let us compute statistics: averages, spreads, and patterns.




Example: Coin flip outcome $\rightarrow$ Number: - Heads $\rightarrow Y = 1$ - Tails $\rightarrow Y = 0$

Discrete vs. Continuous Random Variables

CONTINUOUS RANDOM VARIABLES


- No distinct categories
- No limit on the value
- Tends to be quantitative
- Height
- Weight
- Blood pressure
- Time



Continuous: Area under curve = probability

DISCRETE RANDOM VARIABLES

- Distinct categories (countable)
- No in-between values
- Tends to be qualitative/countable
- Coin flips
- Patient counts at emergency room
- Number of rings in a tree



Discrete: Heights of bars = probability at each point

Continuous: Infinite uncountably many points
→ each single point has $P(Y=y)=0$. Only INTERVALS have non-zero probability.

Discrete: Countably many points → each point can have $P(Y=y)>0$. We sum **probabilities** at individual points.

Discrete Random Variables

- ▶ **Countable** values (often integers)
- ▶ Examples: Coin flips, patient counts at emergency room, number of rings in a tree
- ▶ Use **probability mass function (PMF)** $p(y)$:
 $P(Y = y) = p(y)$ and $\sum_{\text{all } y} p(y) = 1$
- ▶ **Cumulative distribution function (CDF)**:
 $F(y) = P(Y \leq y) = \sum_{x \leq y} p(x)$
- ▶ Also exist quantile functions, moment generating functions, characteristic functions (beyond this course).

Continuous Random Variables

- ▶ Values in an **interval** (uncountably many)
- ▶ Examples: Height, weight, blood pressure, time
- ▶ Use **probability density function (PDF)** $f(y)$
- ▶ Probabilities: $P(a \leq Y \leq b) = \int_a^b f(y) dy$ and $\int_{-\infty}^{\infty} f(y) dy = 1$
- ▶ **Cumulative distribution function (CDF):**
 $F(y) = P(Y \leq y) = \int_{-\infty}^y f(x) dx$


CRITICAL DIFFERENCE: Single Point Probability

CHECKPOINT: Why does $P(Y = 172.5) = 0$ for continuous height?

Property	Discrete	Continuous
$P(Y = y)$	Can be > 0	Always = 0
Example: Fair die	$P(Y = 3) = 1/6$	—
Example: Height	—	$P(\text{height} = 172.5) = 0$
How to find probability	Use $P(Y = y)$	Use $P(a \leq Y \leq b)$

Intuition: Continuous distributions have *uncountably many* points (infinity!), so each single point contributes zero probability.

Only **intervals with area** have non-zero probability.



Part 2: Probability Distributions

Probability Density Function (Continuous)


For a **continuous** Y with density $f(y)$:

$$f(y) \geq 0 \quad \text{for all } y \quad ; \quad \int_{-\infty}^{\infty} f(y) dy = 1$$


For any interval $[a, b]$:

$$P(a \leq Y \leq b) = \int_a^b f(y) dy \quad (\text{shaded area below})$$

Normal: $P(a \leq Y \leq b)$



Chi-Square: $P(a \leq Y \leq b)$



Part 3: Expectation and Variance

Expectation: Discrete Case

For a **discrete** Y with probabilities $p(y)$:

$$E(Y) = \mu = \sum_{\text{all } y} y \cdot p(y)$$

Fair die example ($\{1, 2, 3, 4, 5, 6\}$, each with $p(y) = 1/6$):

$$E(Y) = 1 \cdot \frac{1}{6} + 2 \cdot \frac{1}{6} + \cdots + 6 \cdot \frac{1}{6} = 3.5$$

Interpretation: Long-run average over many rolls.

CHECKPOINT: Understanding $E(Y)$ for the Fair Die

Question: Why is $E(Y) = 3.5$, not 3 or 4? What does 3.5 mean?

- ▶ Is 3.5 a possible outcome when rolling a die? **NO**
- ▶ Does it mean we expect to see 3.5 on the next roll? **NO**
- ▶ Does it mean over 1,000 rolls, the average is *approximately* 3.5? **YES**



Key insight: $E(Y)$ is the **balance point** of the distribution. Imagine the die outcomes as weights on a seesaw—3.5 is where it balances!

Expectation: Continuous Case

For a **continuous** Y with density $f(y)$:

$$E(Y) = \mu = \int_{-\infty}^{\infty} y f(y) dy$$

Interpretation: Center of mass (balance point) of the distribution.



Variance: Discrete Case

For a **discrete** Y :

$$\text{Var}(Y) = \sigma^2 = E((y - \mu)^2) = \sum_{\text{all } y} (y - \mu)^2 \cdot p(y)$$

Fair die with $\mu = 3.5$:



$$\text{Var}(Y) = (1 - 3.5)^2 \cdot \frac{1}{6} + \dots + (6 - 3.5)^2 \cdot \frac{1}{6} \approx 2.92$$

Variance: Continuous Case

For a **continuous** Y :

$$\text{Var}(Y) = \sigma^2 = \int_{-\infty}^{\infty} (y - \mu)^2 f(y) dy$$

Standard deviation: $\sigma = \sqrt{\text{Var}(Y)}$ (same units as Y)



Larger variance = more spread!

CHECKPOINT: Practical Variance Interpretation

Question: Which is less risky: Stock A with $\sigma^2 = 4$ or Stock B with $\sigma^2 = 16$?

- ▶ Stock A: $\sigma = 2$ (tightly concentrated returns)?
- ▶ Stock B: $\sigma = 4$ (widely dispersed returns)?

Key lesson: Variance = **spread of outcomes**. Smaller variance = more predictable!

Part 4: Sample Statistics

Sample Mean and Standard Deviation

Given n observations $y_1, \dots, y_n$:

Sample mean (estimates μ):

$$\bar{y} = \frac{1}{n} \sum_{t=1}^n y_t$$

Sample standard deviation (estimates σ):

$$s = \sqrt{\frac{1}{n-1} \sum_{t=1}^n (y_t - \bar{y})^2}$$

We divide by $(n - 1)$ for unbiased variance estimation (**Bessel's correction**).

Part 5: Rules for Transformations

Linear Transformations: $E(cY)$ and $Var(cY)$

For constant c and random variable Y :

$$E(c) = c,$$

$$E(cY) = c \cdot E(Y)$$

$$Var(c) = 0,$$

$$Var(cY) = c^2 \cdot Var(Y)$$

Key insight: Variance scales with c^2 . Doubling a variable multiplies variance by 4!

CHECKPOINT: can someone say how to prove $\text{Var}(cY) = c^2\text{Var}(Y)$, not $c\text{Var}(Y)$?

Question: If we triple the values (multiply by 3), variance multiplies by 9. Why the square?

Sums of Independent Variables

For independent $Y_1, \dots, Y_n$ with expectations μ_i and variances σ_i^2 and some constants c_i :

Expectation of sum:

$$E \left(\sum_{i=1}^n c_i Y_i \right) = \sum_{i=1}^n c_i \mu_i$$

Variance of sum (independent only):

$$\text{Var} \left(\sum_{i=1}^n c_i Y_i \right) = \sum_{i=1}^n c_i^2 \sigma_i^2$$


Part 6: Key Distributions

The Normal Distribution

Notation: $Y \sim N(\mu, \sigma^2)$

PDF: $f(y) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left\{-\frac{1}{2}\left(\frac{y-\mu}{\sigma}\right)^2\right\}$

Empirical Rule & Visual:



CHECKPOINT: SAT Scores and the 68-95 Rule

Application: SAT scores follow $N(\mu = 500, \sigma = 100)$.

- ▶ What % score between **400-600**? (within 1 standard deviation)
- ▶ What % score between **300-700**? (2 standard deviations away from mean)

CHECKPOINT: SAT Scores and the 68-95 Rule

Application: SAT scores follow $N(\mu = 500, \sigma = 100)$.

- ▶ What % score between **400-600**? (within 1 standard deviation)
 - ▶ **Answer: 68.2%**
- ▶ What % score between **300-700**? (2 standard deviations away from mean)
 - ▶ **Answer: 95.4%**

Why this matters: Standardized tests are designed using normal distributions. Knowing the 68-95 rule helps interpret scores!

Carl Friedrich Gauss (1777–1855)

The “**Prince of Mathematicians**” and namesake of the normal distribution.



Figure 2: German mathematician with groundbreaking contributions to statistics (**error theory** back then), calculus, algebra and number theory, and astronomy

Linear Combinations of Normal Variables

If $Y_1 \sim N(\mu_1, \sigma_1^2)$ and $Y_2 \sim N(\mu_2, \sigma_2^2)$ are *independent*:

$$aY_1 + bY_2 \sim N(a\mu_1 + b\mu_2, a^2\sigma_1^2 + b^2\sigma_2^2)$$


This is **fundamental** for inference about sample means and regression.

The Chi-Square Distribution

Notation: $Y \sim \chi_k^2$ where k = degrees of freedom

Key derivation: $Z^2 \sim \chi_1^2$ if $Z \sim N(0, 1)$; additive property:
 $\sum_{i=1}^k Z_i^2 \sim \chi_k^2$

Properties: $E(Y) = k$, $Var(Y) = 2k$





The Student's t Distribution

Definition: For independent $Z \sim N(0,1)$ and χ_k^2 :

$$T = \frac{Z}{\sqrt{\chi_k^2/k}} \sim t_k$$

t-distribution with small df t-distribution converges to normal



Properties: Symmetric, heavier tails than normal, converges to $N(0,1)$ as $k \rightarrow \infty$

William Sealy Gosset (1876–1937): “Student”



Figure 3: English scientist working as a chemist for Guinness Brewery. Developed the t distribution to handle small samples when variance is unknown (1908)

- ▶ Faced a practical problem: small sample inference for quality control
- ▶ Published under the pseudonym “**Student**” because Guinness forbade employee publications

The t Distribution: Practical Use

Test statistic for population mean:

If $Y_1, \dots, Y_n$ i.i.d. $N(\mu, \sigma^2)$ and $S^2 = \frac{1}{n-1} \sum_{t=1}^n (Y_t - \bar{Y})^2$:

$$T = \frac{\bar{Y} - \mu}{\sqrt{S^2/n}} \sim t_{n-1}$$



This is the **one-sample** t test statistic.

The F Distribution

Notation: $Y \sim F_{u,v}$ where $u, v = \text{degrees of freedom}$

Definition: For independent χ_u^2 and χ_v^2 :

$$F = \frac{\chi_u^2/u}{\chi_v^2/v} \sim F_{u,v}$$



Ronald Aylmer Fisher (1890–1962)



Figure 4: British statistician and geneticist with revolutionary impact on statistics and evolutionary biology. The founder of modern statistics and namesake of the F-distribution.

- ▶ Invented the **F distribution**, **ANOVA** (analysis of variance), and **maximum likelihood estimation**
- ▶ Pioneered **experimental design** principles (randomization, blocking, replication)
- ▶ Won the prestigious **Weldon Medal** for his contributions to statistical science

Legacy: Fisher's work transformed statistics from a descriptive tool to a framework for scientific inference.

Many more distributions of course

