

Multivariate statistics II: PCR and PLSR

Solve Sæbø

Mar 18 2025

Topics for this lecture


- Motivation for dimension reduction in regression
- Principal Components regression (PCR)
- Partial Least Squares Regression (PLSR)

Motivation for dimension reduction in regression

- The ordinary least squares estimator is much used since:
 1. It is an unbiased estimator
 2. Among all unbiased estimators it has minimum variance
- However, it has some problems:
 1. Requires $n > p$ (more observations than variables)
 2. May suffer from inflated variances when $n \sim p$, or...
 3. ...if the predictors are highly collinear

Example - Chemometrics/Spectroscopy

- Near Infra Red (NIR)-spectroscopy absorbance spectra.
- $p = 700$ wavelengths for $n = 100$ samples.



OLS under multi-collinearity, I

- Regression with one wavelength, w1522

```
##
## Call:
## lm(formula = Absorbance ~ NIR.w1522, data = NIRdata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.68013 -0.52481 -0.03636  0.49450  1.56003
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    38.77      7.48    5.184 1.17e-06 ***
## NIR.w1522   -132.94     31.88   -4.169 6.60e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7131 on 98 degrees of freedom
## Multiple R-squared:  0.1507, Adjusted R-squared:  0.142
## F-statistic: 17.38 on 1 and 98 DF,  p-value: 6.598e-05
```

OLS under multi-collinearity, II

- Regression with two highly correlated predictors


```
##
## Call:
## lm(formula = Absorbance ~ NIR.w1522 + NIR.w1532, data = NIRdata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.3721 -0.4781 -0.1079  0.4943  1.5716
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    54.904      8.701   6.310 8.34e-09 ***
## NIR.w1522   -2052.372     592.367  -3.465 0.000792 ***
## NIR.w1532    1878.503     578.968   3.245 0.001614 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6808 on 97 degrees of freedom
## Multiple R-squared:  0.2338, Adjusted R-squared:  0.218
## F-statistic: 14.8 on 2 and 97 DF,  p-value: 2.456e-06
```

A graphical view of multicollinearity

- A dynamic 3D-plot of wavelengths 1522 and 1524 vs protein (Y) using the `rgl` package. Try yourself!

```
library(rgl)
attach(NIRdata)
plot3d(NIR.w1522, NIR.w1532, Absorbance, col= "blue", radius=0.05, type="s", xlab="w1522", ylab="w1532", zlab="Absorbance")
detach(NIRdata)
```

3D plot



OLS and overfitting

- Pick a random set of 100 variables as predictors.
- Fit a linear model
- Compute the R^2 and adjusted R^2

```
vars <- sample(2:701, 50, replace=FALSE)
fmla <- as.formula(paste("Absorbance ~ ", paste(names(NIRdata)[vars], collapse= "+")))
lmfit3 <- lm(fmla, data=NIRdata)
res <- summary(lmfit3)
res$r.squared
res$adj.r.squared
```

```
## [1] 0.930381
```

```
## [1] 0.8593411
```

How to handle these problems

- Forward variable selection in OLS? (data reduction)
- Use alternative estimator, e.g. PCR or PLSR (data reduction to 'principal components')

PCA of NIR data

- Let's return to the 3D plot from above
- Locate the principal components
- Discuss the variances of the two principal components
- How many PC's are needed to capture 99% of the variability in the entire NIR data?

PCR

- In PCR the principal components (PCs) replace the original X-variables as predictors in regression.
- Let a be the number of PCs we choose to use (e.g. explaining 99% X-variability)
- Typically we use $a \ll k$ components in the model.

The model with a components:

$$Y = \alpha_0 + \alpha_1 Z_1 + \alpha_2 Z_2 \cdots + \alpha_a Z_a + \epsilon$$


- The α 's are estimated by the Ordinary Least Squares method.

Example, NIR data

- Run PCA on the Wavelength data, check results, scree plot

```
pcafmla <- as.formula(paste("~ ", paste(colnames(NIRdata)[2:600], collapse= "+")))  
PCAModel.3 <- prcomp(pcafmla, scale.= FALSE, data=NIRdata)  
screeplot(PCAModel.3)
```

The Scree plot



Saving some principal components as extra variables to the data

```
PCAdat <- within(NIRdata, {  
  PC5 <- PCAModel.3$x[,5]  
  PC4 <- PCAModel.3$x[,4]  
  PC3 <- PCAModel.3$x[,3]  
  PC2 <- PCAModel.3$x[,2]  
  PC1 <- PCAModel.3$x[,1]  
})
```

Regression model, PC1-PC5

```
LinearModel.4 <- lm(Y ~ PC1 + PC2 + PC3 + PC4 + PC5, data=PCAdata)
sumfit <- summary(LinearModel.4)
round(sumfit$coefficients,3)
```

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	7.588	0.052	145.807	0.000
## PC1	0.151	0.447	0.337	0.737
## PC2	4.061	1.171	3.468	0.001
## PC3	30.149	3.680	8.192	0.000
## PC4	3.063	4.520	0.678	0.500
## PC5	40.787	6.224	6.553	0.000

- Note that PC1 is not significant


Running the PCR in RStudio with the pls package

- After estimation of the α 's the regression coefficients for the *original* variables may be computed by a reverse transformation.
- These are easiest computed using the `pcr` function in the `pls` package.

```
library(pls)
pcrfit <- pcr(Absorbance ~ ., data=NIRdata, scale=FALSE)
summary(pcrfit)
coefs.pc11 <- coef(pcrfit, 11)
```

The regression coefficients plotted

- Here are the coefficients based on 11 PC's



Better to be presisly wrong...

...than approximately right

The PCR-estimates for the model with original predictors has typically these properties:

- They are biased if $a < k$ (but being closer to zero than OLS)
- Have smaller variances than the least squares estimates.

Prediction error may be smaller for PCR than OLS model since:

Expected squared prediction error =

$$\sigma^2 + \text{var}(\text{estimates}) + \text{bias}(\text{estimates})^2$$

Choosing the number of components, a

- Cross-validation. As Leave-one-out (LOO-CV), or K-fold.
- Root Mean Squared Error of Prediction (RMSEP)

$$RMSEP_{LOO-CV} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_{(i)})^2}$$

where $\hat{y}_{(i)}$ is a prediction of y_i based on a data set where observation i has been left out.

- Compare RMSEP for $a = 1, 2, \dots$ and choose a simple, but good predictor.

LOO-CV example - R-code

```
MVRModel.7 <- pcr(Absorbance ~ ., data=NIRdata, scale=FALSE, ncomp=15, validation="LOO")  
summary(MVRModel.7)
```

Summary with RMSEP-values

Data: X dimension: 100 700

Y dimension: 100 1

Fit method: svdpc

Number of components considered: 15

##

VALIDATION: RMSEP

Cross-validated using 100 leave-one-out segments.

##	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
## CV	0.7737	0.7805	0.7668	0.6406	0.6364	0.5576	0.5386
## adjCV	0.7737	0.7804	0.7667	0.6401	0.6363	0.5575	0.5385
##	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
## CV	0.5499	0.5005	0.4381	0.4328	0.3978	0.4046	0.3979
## adjCV	0.5497	0.5003	0.4375	0.4326	0.3975	0.4044	0.3976
##	14 comps	15 comps					
## CV	0.4041	0.3974					
## adjCV	0.4038	0.3971					


##

TRAINING: % variance explained

##	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps
## X	84.80628	96.599	97.82	98.86	99.44	99.66	99.74
## Absorbance	0.09501	5.334	35.27	37.55	52.14	56.57	56.74
##	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps

Plotting RMSEP

```
plot(RMSEP(MVRModel.7))
```



Partial Least Squares Regression

A problem with PCR is that the PC's may not be good predictors for a response variable.

Partial Least Squares, on the other hand, is seeking PC's that

- have maximum *covariance* between predictors and the response
- are orthogonal to each other
- search more directly for y-relevant information for prediction

An example

Let's sketch a situation where PLS will find other latent components than PCA

Fitting the PLSR model in RStudio

```
MVRModel.8 <- plsr(Absorbance ~ ., data=NIRdata, scale=FALSE, ncomp=15)  
summary(MVRModel.8)
```

Summary

Data: X dimension: 100 700

Y dimension: 100 1

Fit method: kernelpls

Number of components considered: 15


TRAINING: % variance explained

##	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps
## X	45.489	95.56	97.74	98.35	98.76	99.53	99.71
## Absorbance	9.735	17.66	53.15	62.68	72.32	76.26	80.74
##	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps
## X	99.76	99.80	99.85	99.89	99.91	99.94	99.95
## Absorbance	83.11	85.33	86.83	87.75	88.68	89.11	90.35
##	15 comps						
## X	99.97						
## Absorbance	90.71						

Running LOO-CV and RMSEP plot


```
MVRModel.8 <- plsr(Absorbance ~ ., data=NIRdata, scale=FALSE, ncomp=15, validation='LOO')  
plot(RMSEP(MVRModel.8))
```

The RMSEP plot



The regression coefficients plotted

- Here are the coefficients based on 8 PLS PC's



Significance testing in PLSR

- No analytical expression for standard errors of estimates!
- No estimate of noise variance!
- Testing based on something called Jack-knifing

Jack-knifing

- Leave-one-out method
- An estimate of the variance of the regression coefficients is computed from the saved estimated coefficients from the n LOO-estimations.
- Readily computed in the `pls` package using the `jack.test` function.

$$Var(\hat{\beta}) = \frac{n-1}{n} \sum_{i=1}^n (\hat{\beta}_{(-i)} - \hat{\beta})^2$$

Jack-knifing in RStudio

```
MVRModel.8 <- plsr(Absorbance ~ ., data=NIRdata, scale=FALSE, ncomp=15, validation='LOO')
j.test <- jack.test(MVRModel.8, ncomp = 8, use.mean = FALSE)
print(j.test)
```